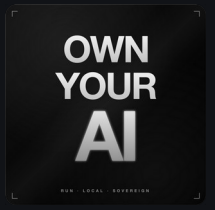
Today’s agenda



Own Your AI • Vienna • 24 September 2026

Intro	**What's new in September 2026**	Andreas Petersson
Focus	**Ten Ways to Hack AI Applications and How to Harden Them**	Andreas Burner
Focus	**(Agentic) Loops At Levels, Live**	Russ Miles
Focus	**Voice AI - Local, Realtime, Best-Practices**	Andi Kunar
Lightning	**Was KI kann**	Florian Zsifkovics
Lightning	**Governance Architecture & Drift Detection**	Jürgen Eckel

News today: ownership • privacy • open weights • tools and hardware

NVIDIA plans to buy Hugging Face

$12.9B

Agreement announced 3 September

Expected close: H1 2027
Regulatory approvals pending

Open, multi-accelerator platform:
NVIDIA's stated commitment.



AI-generated satire

The GPU shop would also own the model bookshop.

NVIDIA announcement and SEC filing, 3 September 2026

Hugging Face is funding local runtimes

February: ggml / llama.cpp team

Historical context: maintainers joined Hugging Face.

22 September: oMLX creator Jun Kim

Hugging Face announces continued Apache 2.0 development.

What matters to users

Maintenance funding, cross-hardware support, and the ability to retain a working stack.

Hugging Face announcements. Existing open-source licenses remain distinct from employment.

“Pacing the frontier”

12 September: Amodei calls for pacing

Commits to embedded outside evaluators. Wider coordination proposed.

22 September: Anthropic ships Opus 5.5

The launch explicitly argues that it fits the pacing commitment.

Safety policy and competitive incentives

“Marketing stunt” is a critical interpretation, not an established motive.

Amodei essay, Anthropic launch. Contemporary criticism is attributed in notes.

Concrete frontier-AI rules

September brings concrete oversight mechanisms and new proposals

California: laws signed

Assessor / auditor framework.
Executive order requests kill-switch advice.

US Senate: proposal, 23 Sep

Frontier pre-certification and pause powers.

UK: first reading, 8 Sep

Private member’s superintelligence bill.



“From my cold dead hands”

No blanket ban on owning a GPU in these measures.

California: 9 and 18 September. US Senate: 23 September. UK: 8 September.

ZCode: what leaves the workspace?

Researchers found background snapshots

Uploads went beyond the expected conversation context.

The control gap mattered

The tested local toggles did not stop the collection path.

The vendor removed the workflow

v3.14.0 fixed it. A researcher checked v3.14.3 on 23 September.

Researcher reports, 18 and 23 September. Official changelog and vendor response.

Muse: a lock is only the beginning

Meta’s personal agent began its US rollout on 8 September

Today: Secure VM

Dedicated cloud computer and browser.
Policies limit Meta access; access is possible.

Planned: Confidential VM

User-held key and operator exclusion.
Announced for later in 2026.



The cables still need explaining.

Tenant isolation, agent permissions and operator privacy are different claims.

Meta launch and technical account, 8 September 2026

What a confidential VM must prove

Generic TEE explanation. Muse’s exact implementation remains unverified.

Hardware protects the running workload

Memory isolation narrows what the host administrator can inspect.

Attestation identifies measured software state

A verifier checks measurements and security versions before releasing secrets.

The protected program still sees usable data

Its permissions, bugs, and outbound paths still matter.

General background: Intel TDX documentation. This does not identify Muse’s hardware.

Muse’s privacy boundary

Current published behavior	Question for confidential mode
Inference and telemetry leave the VM.	Does their processing stay within an attested confidential boundary?
Training uses sanitized trajectories unless the user opts out.	Which logging, training and debugging paths will the private mode permit?
Sentinel controls network access.	What does each authorized external service receive?

A lock around the VM is only part of an end-to-end privacy argument.

Meta technical account. Right column: our verification questions, not claimed defects.

Muse: evidence to look for at launch

Optional depth: a verification agenda

Boundary	Evidence needed
Keys and recovery	Device loss, backup access, deletion and restore rules
Attestation and updates	Accepted measurements, signed builds, rollback policy
Inference and telemetry	Every plaintext hop, logs and safety-monitoring paths
Independent review	Public audit findings and build-to-source correspondence

These are open questions, not findings of a broken implementation.

Our analysis. No public complete Muse confidential specification verified at cutoff.

Who receives the request?

Anthropic’s September report alleges relaying

Selected user requests reached another provider without clear visibility.

The claims differ across companies

Moonshot / DeepSeek: served responses. Xiaomi: replayed saved sessions.

The local engineering question

Can we inspect the complete processing path and enforce its boundaries?

Attributed allegations from Anthropic. Independent verification not established here.

MiMo V2.6: yes, the big weights are open

Xiaomi • released 22 September • MIT declared in all three model cards

Pro-RL • 1.02T total / 42B active

573GB published weights. Server scale.

Flash-RL • 309B total / 15B active

178GB weights. Still a very large model.

Distill-Qwen-9B • 9.41B total

19GB BF16. Qwen3.5-9B + MiMo-data SFT.

Pro + Flash: 1M context, text / image / video / audio.
Payload sizes are not runtime memory requirements.



Active parameters do not shrink the luggage.

Official Xiaomi release, Hugging Face model cards and payload inventory. Checked 24 Sep.

Four small-model stories, four different origins

A compact comparison of where the weights came from and what they are for

Model / size	Creator and origin	Use / boundary
K2 Horizon 7B 7B core / 9B total	IFM / MBZUAI Own pretraining + specialist RL	General model, 512K context. Separate from Kimi. Apache 2.0.
Needle 3 29M to 121M	Cactus Compute Purpose-built tool-calling network	8 to 29MB files; tools / extraction. Narrow scope; review telemetry.
MiMo Distill Qwen-9B 9.41B total	Xiaomi: Qwen3.5-9B base SFT on MiMo-generated data	Agent-tuning starting point. MIT declared; not Pro in 9B.
MiniCPM5-2B 2.517B total	OpenBMB: own model pipeline Distills 16 specialist RL experts	General / agent tasks; GGUF, MLX. Apache 2.0; 128K cache adds RAM.

Primary model cards and training recipes. Sizes include embeddings where reported.

Bonsai 2: 27B in a smaller pot

5.95 GB

PrismML ternary compression
of Alibaba Qwen3.8-27B



Same 27B parameters, smaller storage.
Custom Prism runtime; cache needs extra memory.

PrismML, 17 September. AI-generated conceptual illustration.

DeepSeek V4.1 Flash is server-scale

16B

Active parameters
per output token

552B backbone
+ 196B Engram

Native vision
1M context

Active parameters describe work per token. Total weights drive storage and memory.

DeepSeek release and model card, 10 September 2026

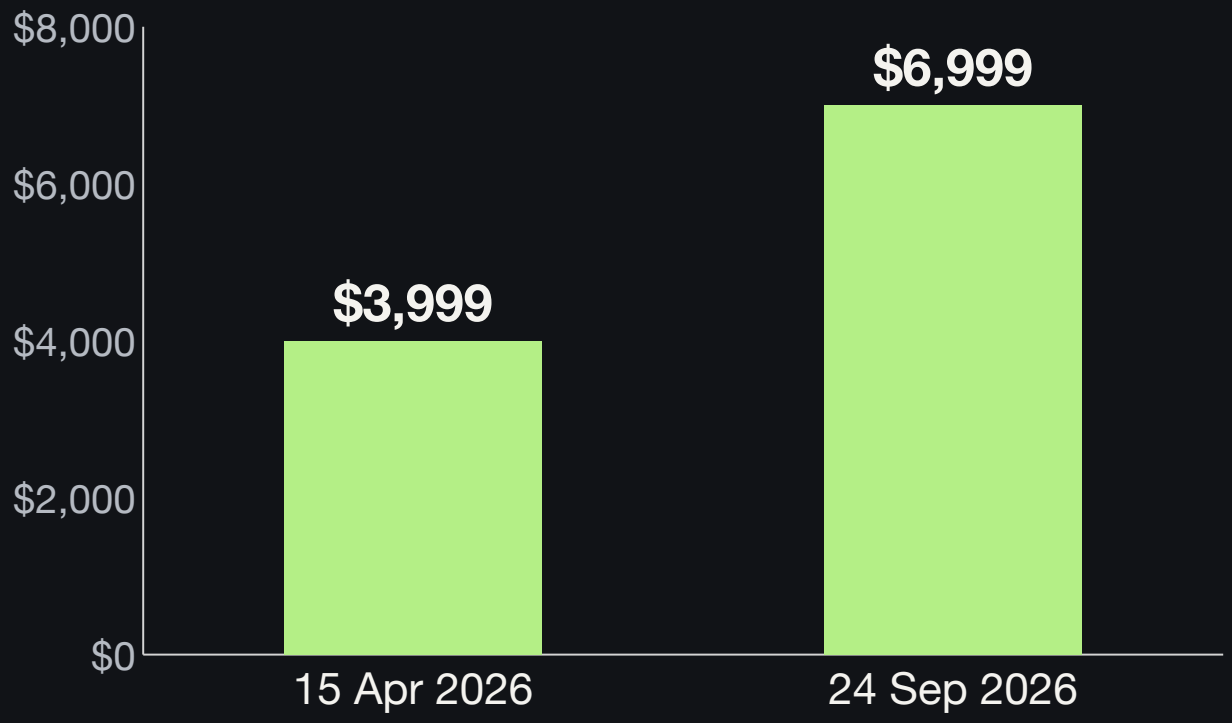
Voice and images: who built what?

Release	Creator and inherited model	Local angle / boundary
VibeVoice Streaming ASR	Microsoft Research: offline VibeVoice-ASR with Qwen2.5 + audio encoders	7B backbone, 8.67B total. Streaming; German. MIT.
Parakeet Redux 22 September	Moondream ternary conversion of NVIDIA Parakeet TDT 0.6B v3	178MB; 25 languages. CC BY 4.0; test noise / German.
Qwen-Image-2.1	Alibaba Qwen team Image generation and editing	Local image model. Commercial license required.

Microsoft, Moondream and Qwen primary sources

My AI box has entered its Bitcoin era

ASUS Ascent GX10, same 2TB configuration, US prices before sales tax





+75%

NVIDIA DGX Spark US store:
“Out of Stock”

Advertised prices, not resale returns. The increase spans April to September.

ASUS article and store, NVIDIA store. Stock snapshot: 24 September 2026

Your Mac now ships a local AI command

Official Apple fm CLI • macOS 27 became available on 14 September

fm respond "Explain why my AI box needs a mortgage."

Built into macOS 27

fm respond for scripts. fm chat for interactive use.

On-device Foundation Models

Compatible Apple Silicon + Apple Intelligence required.

OS access, not downloadable open weights

Apple’s AFM 3 family was custom-built with Google.

The distribution story: a local model is now one terminal command away.

Apple WWDC26 CLI session, September OS release and AFM 3 research. Example prompt, no output shown.

PAIR turns LAN machines into workers

NVIDIA beta announced 3 September

Concurrent requests across machines

RTX PCs, DGX Spark and Apple M4+ Macs can share the workload.

Ollama and LM Studio integration

Useful for several local jobs or users at once.

Each request still needs a capable worker

PAIR does not pool VRAM to fit a single larger model.

NVIDIA IFA announcement. RTX Spark Windows systems are scheduled for October.

The local runtime update

Tool	September change	Before trying it
Splash	Specialized Apple inference with DFlash 2	M3+, macOS 26.4+, 36GB minimum.
Ollama 0.34	MLX support matured. 0.34.4 landed 23 Sep.	Check model support and memory behavior.
llama.cpp 0.5	23 Sep release: kernels and server fixes.	Gains depend on backend, model and workload.

Update one component at a time and keep a reproducible local comparison.

Project release notes and Inco / LM Studio announcements

Jev / Laya: decisions without a chatty model

Jev

TypeSafe hosted decision family. Base weights undisclosed.

Laya English · 421M

Convai Innovations: ModernBERT-large + RL decision head.

Laya multilingual · 322M

Convai Innovations: mmBERT-base + RL decision head.

“Customer reports a duplicate charge.”

billing / technical / other → choice + score

laya-mlx is an independent Apple port. Typed decisions still need permission checks.

Jev: TypeSafe, 15 September. Laya: model repository created 18 September.

Laya: comparable numbers need context

Independent laya-mlx benchmark, M3 Max 40-core GPU, 128GiB

Checkpoint	MPS FP32	MLX FP32	MLX FP16
English	22.70 ms	18.55 ms	17.75 ms
Multilingual	13.60 ms	10.73 ms	10.91 ms

Short input, one question, warm runs. Model loading and downloads are excluded.

Pinned BENCHMARKS.md revision, 22 September. 5 warmups, 50 measured iterations.

Local agents are becoming products

Perplexity Portable Computer on Windows

14 September. NVIDIA GPUs with at least 24GB VRAM.

Local work, explicit cloud escalation

Vendor says cloud use needs permission. Local tasks use no Perplexity credits.

Cost and privacy still depend on the whole path

Electricity and hardware cost remain. Websites and connectors still receive authorized data.

NVIDIA / Perplexity launch material. Local tasks are not a promise of universal offline behavior.

The laptop memory ceiling is moving

HP / AMD ZBook Ultra G3a announced 15 September

192GB

Maximum unified memory

160GB

Maximum GPU allocation

Expected in October. Price and independent workload results still to come.

HP announcement. Configuration maximums, not a measured performance claim.

Local Qwen image generation

Built-in image tool





Local Qwen Image 2.1 Q8





Local prompt excerpts

GPU ownership

“A lovable mustachioed cowboy in a dusty frontier town at sunset hugs two enormous modern graphics cards with clearly recognizable cooling fans and circuit boards as his most treasured possessions.”

Hardware prices

“a small gold cube-shaped personal AI computer wearing an oversized astronaut helmet is riding a tiny homemade toy rocket toward a cartoon cheese moon.”

1152×768, 40 Euler steps, CFG 6. M1 Ultra, 48 GPU cores, 128 GiB. Local: 17.5 / 17.0 min.

Similar prompts, not a controlled benchmark. Full local prompts and seeds in speaker notes.

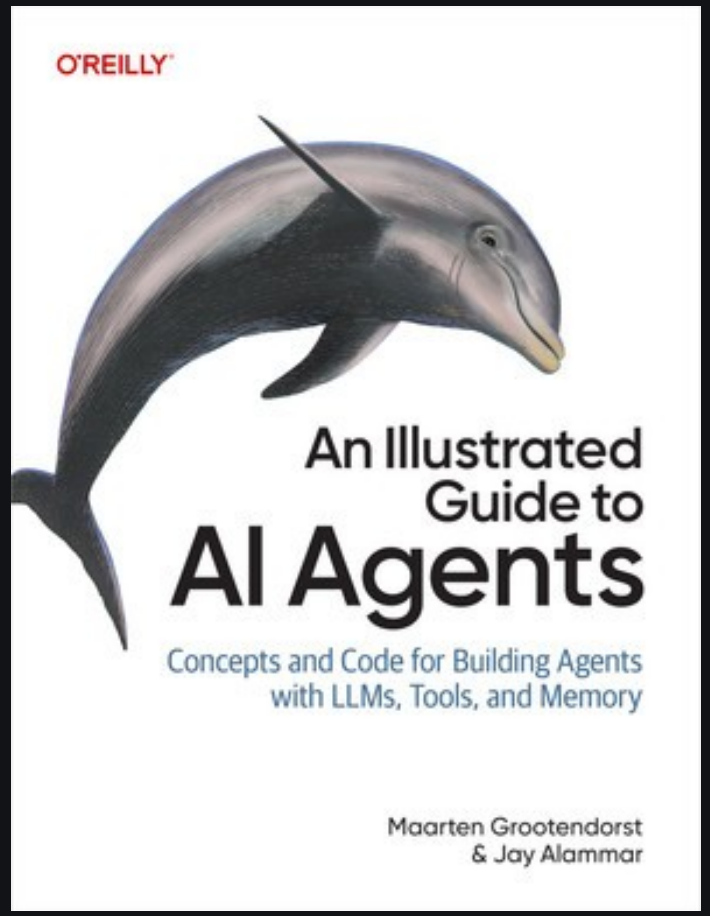
AGNTCon + MCPCon: agent infrastructure

Amsterdam • 17 and 18 September 2026 • Conference takeaways

Theme	What it enables	Status / boundary
Out-of-the-box software factories	Reusable coding pipelines: agents, environments and review controls.	Products exist; quality still needs validation.
A2A 1.0	Agents discover peers and delegate work. Version negotiation; signed agent cards.	Stable since March. Conference context.
MCP Stateless	Each request carries its own context. Simpler scaling and load balancing.	2026-07-28 specification. Check client / SDK support.
WebMCP	Websites expose tools through JS / HTML. Actions use the current browser context.	Chrome 149 origin trial. Support is still evolving.

Conference programme; official A2A, MCP and Chrome documentation. Checked 24 September.

Book of the month



Andreas' recommendation

An Illustrated Guide to AI Agents

Maarten Grootendorst & Jay Alammar

Understand how an agent is put together

Build the TinyAgent as you read

Learn to assess results and reliability

oreilly.com • ISBN 9798341662681

O'Reilly • September 2026 • 450 pages

Stay connected. Up next…

Join us on Discord

discord.gg/J2XzvVND66

Next on today's programme

Ten Ways to Hack AI Applications and How to Harden Them
Andreas Burner

(Agentic) Loops At Levels, Live
Russ Miles

Voice AI - Local, Realtime, Best-Practices
Andi Kunar

Was KI kann
Florian Zsifkovics • Lightning talk

Governance Architecture & Drift Detection

Jürgen Eckel • Lightning talk